

# lung\_cancer

June 8, 2022

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
```

```
[ ]: # read the dataset or import the data using pandas
# 1 -- > NO
# 2 -- > YES

df = pd.read_csv('C:/Users/91861/Downloads/survey lung cancer.csv')
df.head()
```

```
[ ]:  GENDER  AGE  SMOKING  YELLOW_FINGERS  ANXIETY  PEER_PRESSURE  \
0      M    69      1           2           2           1
1      M    74      2           1           1           1
2      F    59      1           1           1           2
3      M    63      2           2           2           1
4      F    63      1           2           1           1

      CHRONIC DISEASE  FATIGUE  ALLERGY  WHEEZING  ALCOHOL CONSUMING  COUGHING  \
0                   1         2         1         2           2         2
1                   2         2         2         1           1         1
2                   1         2         1         2           1         2
3                   1         1         1         1           2         1
4                   1         1         1         2           1         2

      SHORTNESS OF BREATH  SWALLOWING DIFFICULTY  CHEST PAIN  LUNG_CANCER  \
0                        2                      2           2        YES
1                        2                      2           2        YES
2                        2                      1           2        NO
3                        1                      2           2        NO
4                        2                      1           1        NO
```

Unnamed: 16

```

0      NaN
1      NaN
2      NaN
3      NaN
4      NaN

```

```
[ ]: df.drop(columns=df.columns[-1],
            axis=1,
            inplace=True)
```

```
[ ]: df.head()
```

```
[ ]:
  GENDER  AGE  SMOKING  YELLOW_FINGERS  ANXIETY  PEER_PRESSURE  \
0      M   69        1                2        2              1
1      M   74        2                1        1              1
2      F   59        1                1        1              2
3      M   63        2                2        2              1
4      F   63        1                2        1              1

  CHRONIC DISEASE  FATIGUE  ALLERGY  WHEEZING  ALCOHOL CONSUMING  COUGHING  \
0                1        2        1        2                2        2
1                2        2        2        1                1        1
2                1        2        1        2                1        2
3                1        1        1        1                2        1
4                1        1        1        2                1        2

  SHORTNESS OF BREATH  SWALLOWING DIFFICULTY  CHEST PAIN  LUNG_CANCER
0                    2                      2          2        YES
1                    2                      2          2        YES
2                    2                      1          2        NO
3                    1                      2          2        NO
4                    2                      1          1        NO

```

```
[ ]: # Exploratory Data Analysis
# Now, let's see the size of the dataset

df.shape
```

```
[ ]: (309, 16)
```

```
[ ]: df.info()

# Out of 16 features, we have 14 int types and only two with the string data_
↳ types.
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 309 entries, 0 to 308
```

Data columns (total 16 columns):

#	Column	Non-Null Count	Dtype
0	GENDER	309 non-null	object
1	AGE	309 non-null	int64
2	SMOKING	309 non-null	int64
3	YELLOW_FINGERS	309 non-null	int64
4	ANXIETY	309 non-null	int64
5	PEER_PRESSURE	309 non-null	int64
6	CHRONIC DISEASE	309 non-null	int64
7	FATIGUE	309 non-null	int64
8	ALLERGY	309 non-null	int64
9	WHEEZING	309 non-null	int64
10	ALCOHOL CONSUMING	309 non-null	int64
11	COUGHING	309 non-null	int64
12	SHORTNESS OF BREATH	309 non-null	int64
13	SWALLOWING DIFFICULTY	309 non-null	int64
14	CHEST PAIN	309 non-null	int64
15	LUNG_CANCER	309 non-null	object

dtypes: int64(14), object(2)

memory usage: 38.8+ KB

```
[ ]: df.isnull().sum()

# Woah! Fortunately, this dataset doesn't hold any missing values.
```

```
[ ]: GENDER      0
      AGE        0
      SMOKING    0
      YELLOW_FINGERS  0
      ANXIETY    0
      PEER_PRESSURE  0
      CHRONIC DISEASE  0
      FATIGUE    0
      ALLERGY    0
      WHEEZING   0
      ALCOHOL CONSUMING  0
      COUGHING   0
      SHORTNESS OF BREATH  0
      SWALLOWING DIFFICULTY  0
      CHEST PAIN  0
      LUNG_CANCER  0
      dtype: int64
```

```
[ ]: df.describe()
```

```
# As we are getting some information from each feature so let's see how
↳ statistically the dataset is spread
```

```
[ ]:
```

	AGE	SMOKING	YELLOW_FINGERS	ANXIETY	PEER_PRESSURE	\
count	309.000000	309.000000	309.000000	309.000000	309.000000	
mean	62.673139	1.563107	1.569579	1.498382	1.501618	
std	8.210301	0.496806	0.495938	0.500808	0.500808	
min	21.000000	1.000000	1.000000	1.000000	1.000000	
25%	57.000000	1.000000	1.000000	1.000000	1.000000	
50%	62.000000	2.000000	2.000000	1.000000	2.000000	
75%	69.000000	2.000000	2.000000	2.000000	2.000000	
max	87.000000	2.000000	2.000000	2.000000	2.000000	

	CHRONIC_DISEASE	FATIGUE	ALLERGY	WHEEZING	ALCOHOL_CONSUMING	\
count	309.000000	309.000000	309.000000	309.000000	309.000000	
mean	1.504854	1.673139	1.556634	1.556634	1.556634	
std	0.500787	0.469827	0.497588	0.497588	0.497588	
min	1.000000	1.000000	1.000000	1.000000	1.000000	
25%	1.000000	1.000000	1.000000	1.000000	1.000000	
50%	2.000000	2.000000	2.000000	2.000000	2.000000	
75%	2.000000	2.000000	2.000000	2.000000	2.000000	
max	2.000000	2.000000	2.000000	2.000000	2.000000	

	COUGHING	SHORTNESS_OF_BREATH	SWALLOWING_DIFFICULTY	CHEST_PAIN
count	309.000000	309.000000	309.000000	309.000000
mean	1.579288	1.640777	1.469256	1.556634
std	0.494474	0.480551	0.499863	0.497588
min	1.000000	1.000000	1.000000	1.000000
25%	1.000000	1.000000	1.000000	1.000000
50%	2.000000	2.000000	1.000000	2.000000
75%	2.000000	2.000000	2.000000	2.000000
max	2.000000	2.000000	2.000000	2.000000

```
[ ]: df.corr().head()

# It is always better to check the correlation between the features,
# so that we can analyze that which feature is negatively correlated and which
↳ is positively correlated
# Let's check the correlation between various features.
```

```
[ ]:
```

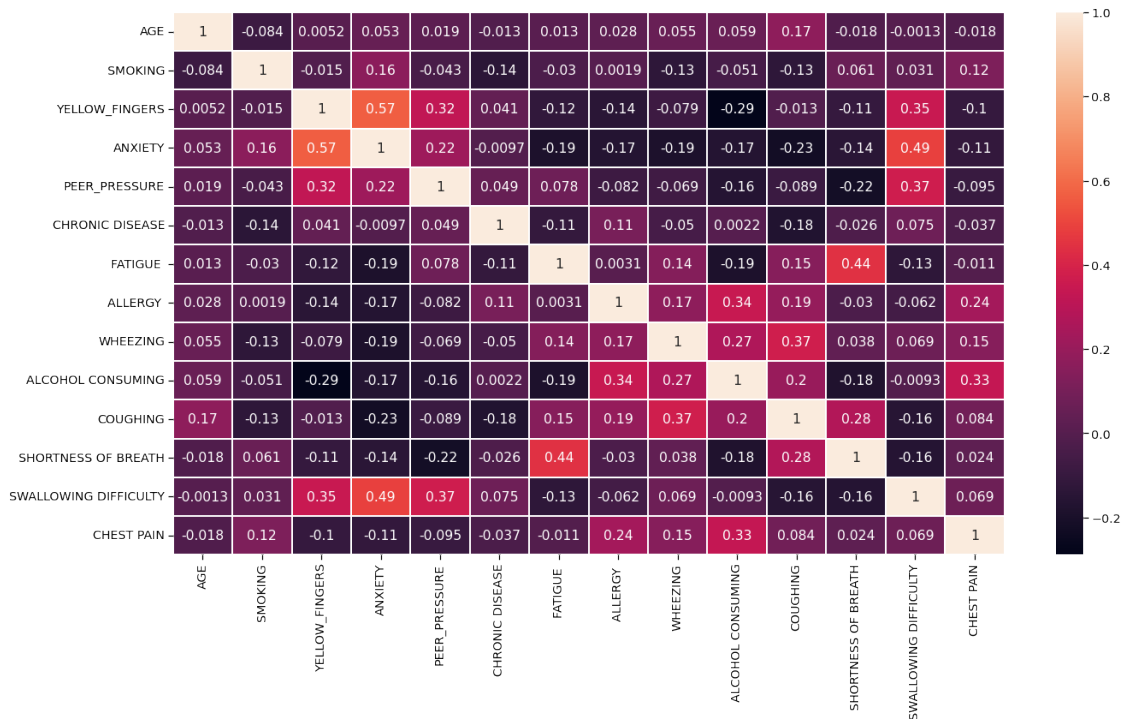
	AGE	SMOKING	YELLOW_FINGERS	ANXIETY	PEER_PRESSURE	\
AGE	1.000000	-0.084475	0.005205	0.053170	0.018685	
SMOKING	-0.084475	1.000000	-0.014585	0.160267	-0.042822	
YELLOW_FINGERS	0.005205	-0.014585	1.000000	0.565829	0.323083	
ANXIETY	0.053170	0.160267	0.565829	1.000000	0.216841	
PEER_PRESSURE	0.018685	-0.042822	0.323083	0.216841	1.000000	

	CHRONIC DISEASE	FATIGUE	ALLERGY	WHEEZING	\
AGE	-0.012642	0.012614	0.027990	0.055011	
SMOKING	-0.141522	-0.029575	0.001913	-0.129426	
YELLOW_FINGERS	0.041122	-0.118058	-0.144300	-0.078515	
ANXIETY	-0.009678	-0.188538	-0.165750	-0.191807	
PEER_PRESSURE	0.048515	0.078148	-0.081800	-0.068771	

	ALCOHOL CONSUMING	COUGHING	SHORTNESS OF BREATH	\
AGE	0.058985	0.169950	-0.017513	
SMOKING	-0.050623	-0.129471	0.061264	
YELLOW_FINGERS	-0.289025	-0.012640	-0.105944	
ANXIETY	-0.165750	-0.225644	-0.144077	
PEER_PRESSURE	-0.159973	-0.089019	-0.220175	

	SWALLOWING DIFFICULTY	CHEST PAIN
AGE	-0.001270	-0.018104
SMOKING	0.030718	0.120117
YELLOW_FINGERS	0.345904	-0.104829
ANXIETY	0.489403	-0.113634
PEER_PRESSURE	0.366590	-0.094828

```
[ ]: plt.figure(figsize=(20,12))
sns.set_context('notebook',font_scale = 1.3)
sns.heatmap(df.corr(),annot=True,linewidth =2)
plt.tight_layout()
```



```
[ ]: # Getting Started

# Title : Lung Cancer Prediction

# Lung Cancer Status :


# 0 -- > Yes

# 1 -- > NO

df.replace({"LUNG_CANCER":{"YES":0,'NO':1}},inplace=True)
```

```
[ ]: # Age Analysis

plt.figure(figsize=(15,6))
sns.set_context('notebook',font_scale = 1.5)
sns.barplot(x=df.AGE.value_counts().index,y=df.AGE.value_counts().values)
plt.tight_layout()
```



```
[ ]: df['AGE'].value_counts()
```

```
[ ]: 64      20
      63      19
      56      19
      62      18
      60      17
      61      16
      59      15
      70      15
      67      13
```

```
58    13
69    11
55    11
72    10
71    10
68     9
57     9
77     9
51     8
54     8
65     7
74     6
75     5
76     4
52     4
53     4
73     4
47     4
66     4
49     3
81     2
78     2
44     2
48     2
21     1
79     1
38     1
39     1
87     1
46     1
Name: AGE, dtype: int64
```

```
[ ]: minAge=min(df.AGE)
maxAge=max(df.AGE)
meanAge=df.AGE.mean()
print('Min Age :',minAge)
print('Max Age :',maxAge)
print('Mean Age :',meanAge)
print()
```

```
Min Age : 21
Max Age : 87
Mean Age : 62.67313915857605
```

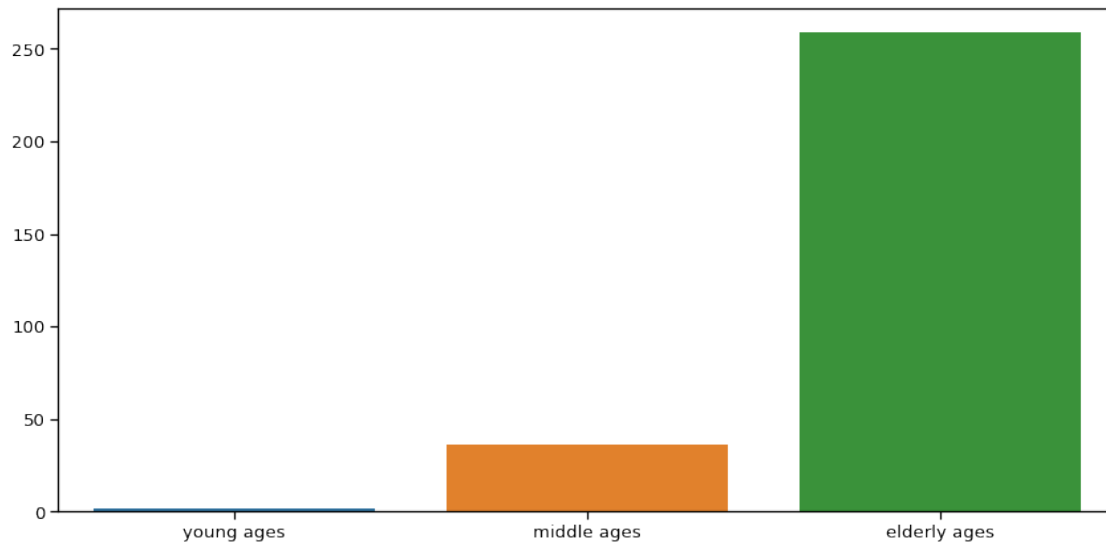
```
[ ]: # We should divide the Age feature into three parts - "Young", "Middle" and
      ↪ "Elder"
```

```

Young = df[(df.AGE>=29)&(df.AGE<40)]
Middle = df[(df.AGE>=40)&(df.AGE<55)]
Elder = df[(df.AGE>55)]

plt.figure(figsize=(12,6))
sns.set_context('notebook',font_scale = 1.2)
sns.barplot(x=['young ages','middle ages','elderly_
↪ages'],y=[len(Young),len(Middle),len(Elder)])
plt.tight_layout()

```



```

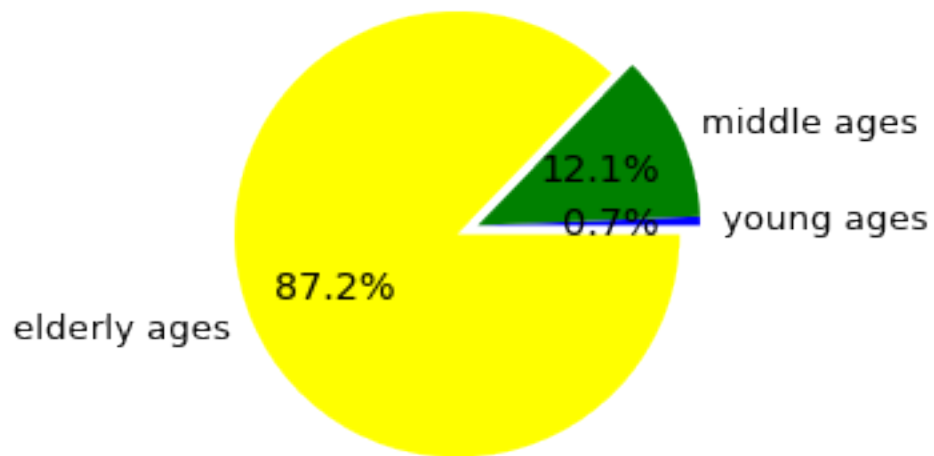
[ ]: # Inference: Here we can see that elder people are the most affected by heart_
↪disease and young ones are the least affected.

# To prove the above inference we will plot the pie chart.

colors = ['blue','green','yellow']
explode = [0,0,0.1]
plt.figure(figsize=(5,5))
sns.set_context('notebook',font_scale = 1.2)
plt.pie([len(Young),len(Middle),len(Elder)],labels=['young ages','middle_
↪ages','elderly ages'],explode=explode,colors=colors, autopct='%1.1f%%')
plt.tight_layout()

```

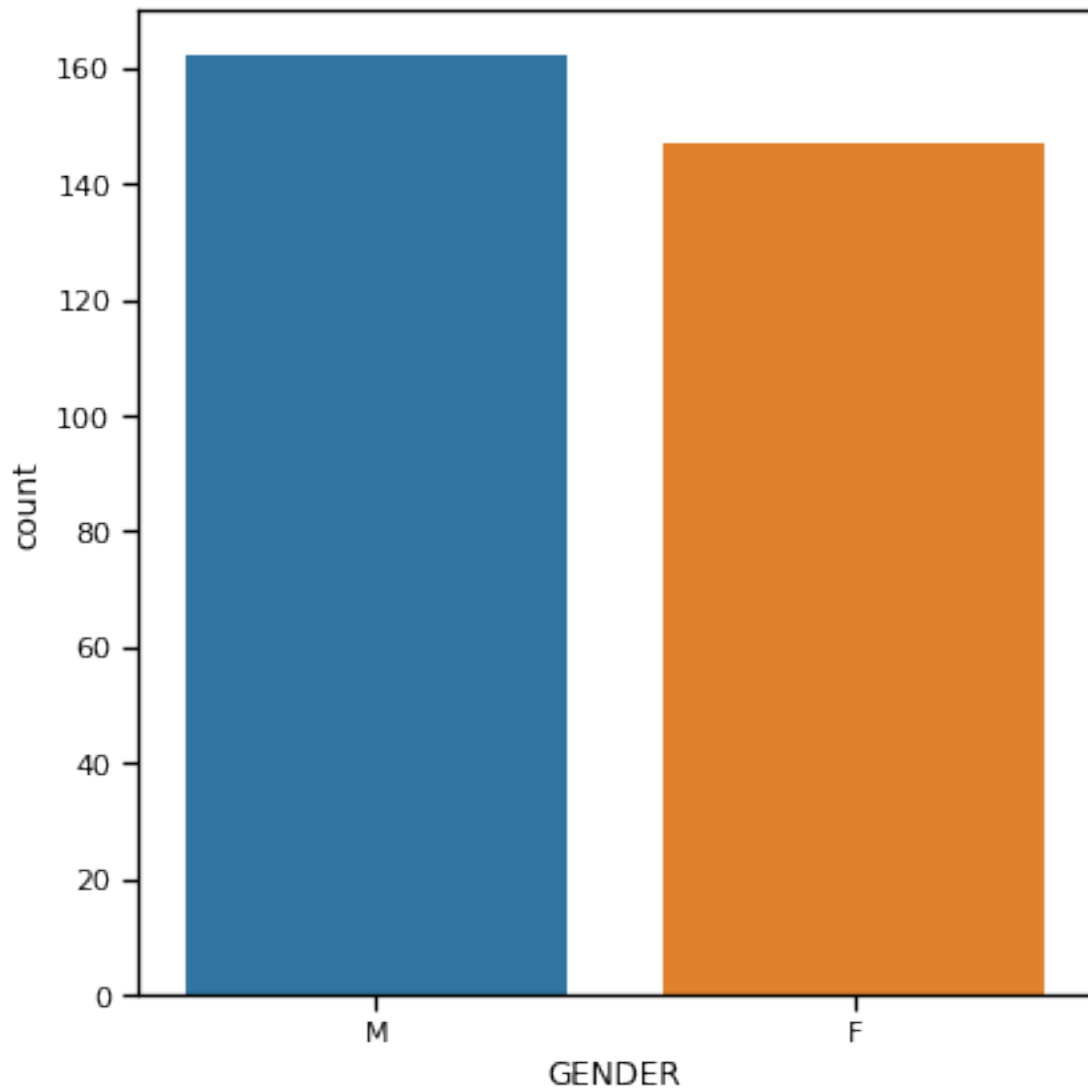




```
[ ]: # Sex - Feature Analysis
```

```
plt.figure(figsize=(6,6))  
sns.set_context('notebook',font_scale = 1)  
sns.countplot(df['GENDER'])  
plt.tight_layout()
```


```
# Inference: Here it is clearly visible that, Ratio of Male to Female almost  
↪ equal but male is higher.
```



```
[ ]: # Lung cancer - Feature Analysis
      # Most affected cancer


      colours=["#f7b2b0", "#8f7198", "#003f5c"]
      sns.countplot(data= df, x="LUNG_CANCER", palette=colours)
```

```
[ ]: <AxesSubplot:xlabel='LUNG_CANCER', ylabel='count'>
```



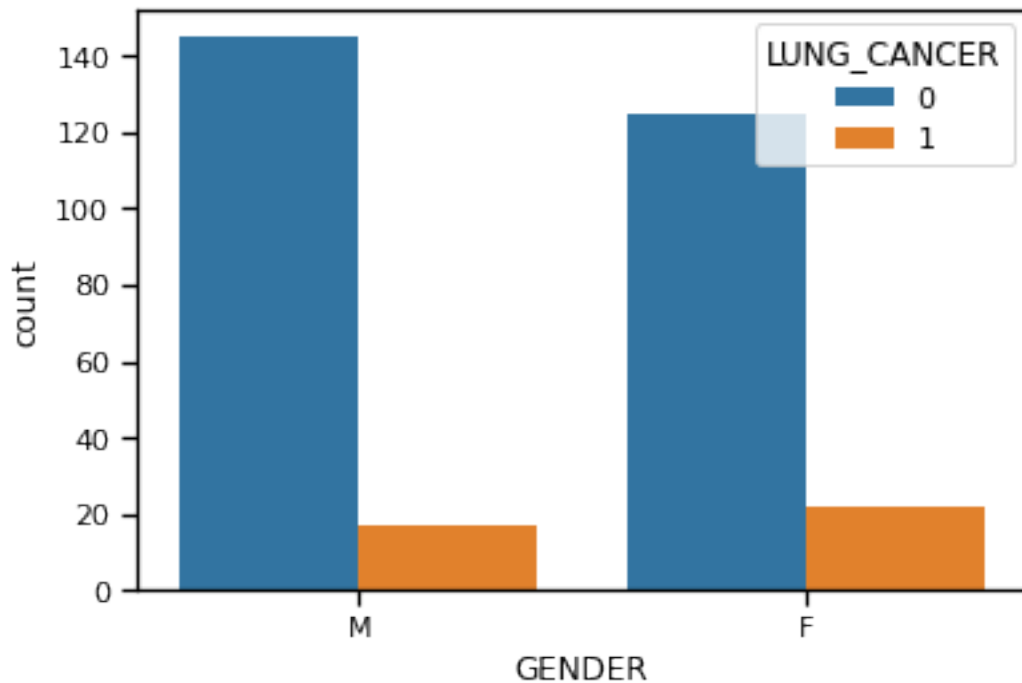
```
[ ]: sns.countplot(x='LUNG_CANCER',hue='LUNG_CANCER',data=df)
```

```
[ ]: <AxesSubplot:xlabel='LUNG_CANCER', ylabel='count'>
```



```
[ ]: sns.countplot(x='GENDER',hue='LUNG_CANCER',data=df)
```

```
[ ]: <AxesSubplot:xlabel='GENDER', ylabel='count'>
```

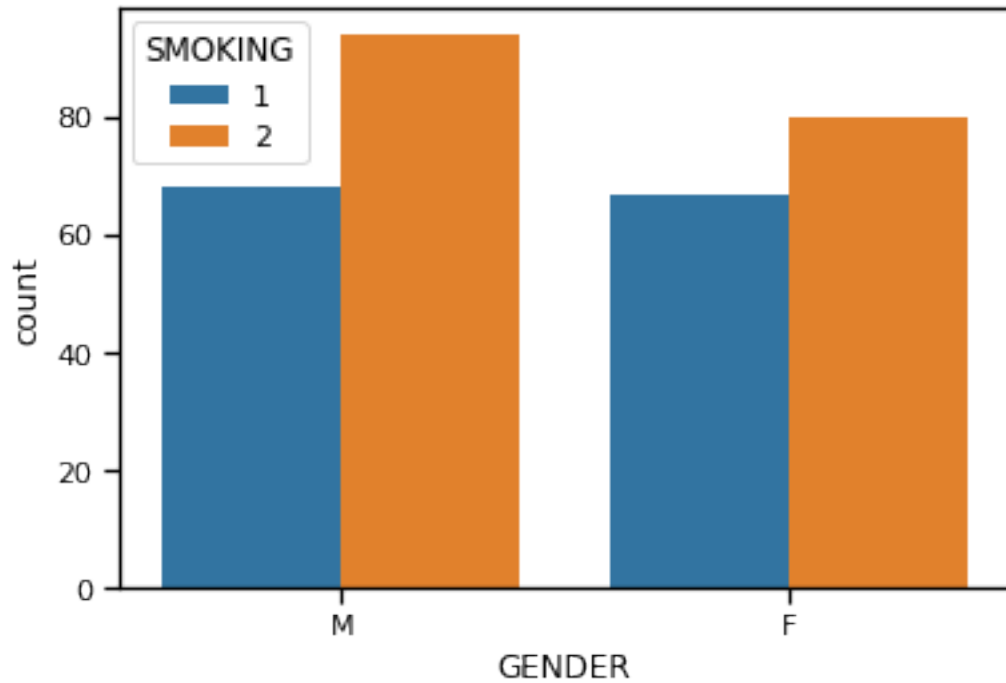


```
[ ]: sns.countplot(x='GENDER',hue='SMOKING',data=df)
```

```
df['SMOKING'].value_counts()
```

```
# MALE SMOKING MORE THAN FEMALE
```

```
[ ]: 2    174  
     1    135  
     Name: SMOKING, dtype: int64
```



```
[ ]: df['LUNG_CANCER'].value_counts()
```

```
[ ]: 0    270
      1     39
      Name: LUNG_CANCER, dtype: int64
```

```
[ ]: # encode the Gender - Feature Analysis

df.replace({"GENDER":{"M":0,'F':1}},inplace=True)
df.head(5)
```

```
[ ]:   GENDER  AGE  SMOKING  YELLOW_FINGERS  ANXIETY  PEER_PRESSURE  \
0      0   69      1           2           2           1
1      0   74      2           1           1           1
2      1   59      1           1           1           2
3      0   63      2           2           2           1
4      1   63      1           2           1           1

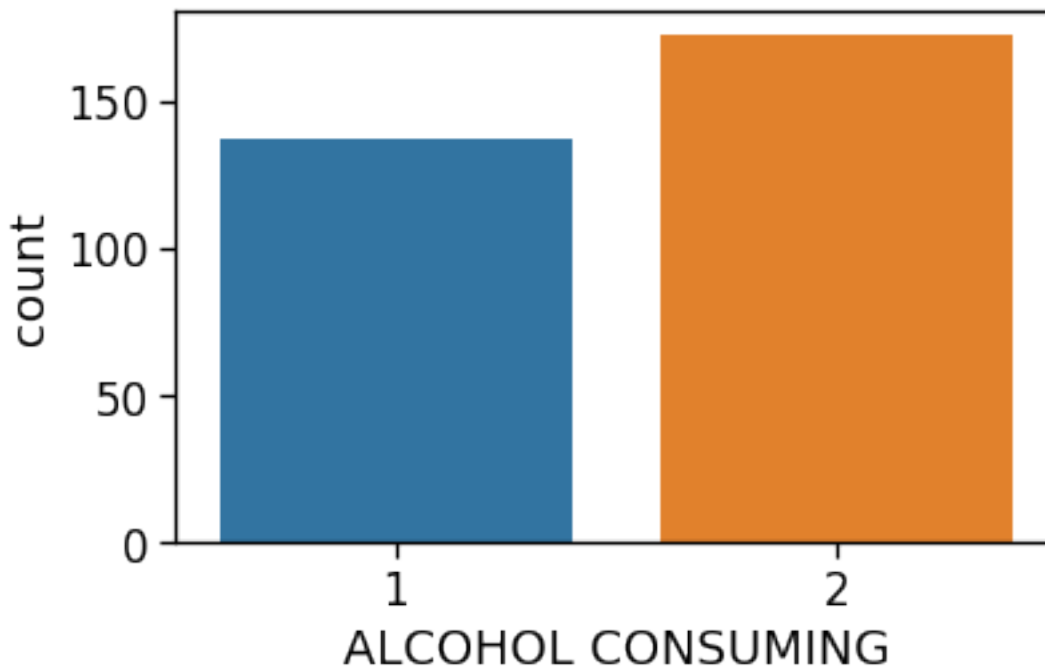
   CHRONIC_DISEASE  FATIGUE  ALLERGY  WHEEZING  ALCOHOL_CONSUMING  COUGHING  \
0                1        2        1        2                2        2
1                2        2        2        1                1        1
2                1        2        1        2                1        2
3                1        1        1        1                2        1
4                1        1        1        2                1        2
```

	SHORTNESS OF BREATH	SWALLOWING DIFFICULTY	CHEST PAIN	LUNG_CANCER
0	2	2	2	0
1	2	2	2	0
2	2	1	2	1
3	1	2	2	1
4	2	1	1	1

```
[ ]: df['GENDER']=df['GENDER'].astype(np.int64)
```

```
[ ]: # Alcohol consuming - Feature Analysis

sns.set_context('notebook',font_scale = 1.5)
sns.countplot(df['ALCOHOL CONSUMING'])
plt.tight_layout()
```



```
[ ]: # TARGET VARIABLE IS LUNG CANCER

X = df.drop(columns=['LUNG_CANCER'],axis=1)
y = df['LUNG_CANCER']
```

```
[ ]: print("The shape of X is " , X.shape)
print("The shape of Y is " , y.shape)
```

The shape of X is (309, 15)

The shape of Y is (309,)

```
[ ]: from sklearn.model_selection import train_test_split

# separating into train and testing

X_train, X_test, Y_train, Y_test = train_test_split(X, y, test_size=0.
↳ 2, stratify=y, random_state=42)

print("Shape of X_train is ", X_train.shape)
print("Shape of X_test is ", X_test.shape)
print("Shape of Y_train is ", Y_train.shape)
print("Shape of Y_test is ", Y_test.shape)
```

Shape of X\_train is (247, 15)  
Shape of X\_test is (62, 15)  
Shape of Y\_train is (247,)  
Shape of Y\_test is (62,)

```
[ ]: print(Y_train.value_counts())
print(Y_test.value_counts())
```

```
0    216
1     31
Name: LUNG_CANCER, dtype: int64
0     54
1      8
Name: LUNG_CANCER, dtype: int64
```

```
[ ]: from sklearn import preprocessing
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.svm import LinearSVC
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import cross_val_score
from sklearn.metrics import precision_score, recall_score, confusion_matrix,
↳ classification_report, accuracy_score, f1_score
from sklearn import metrics
from sklearn.neighbors import KNeighborsClassifier
```

```
[ ]: #pipelines of models( it is short was to fit and pred)

pipeline_lr=Pipeline([('lr_classifier', LogisticRegression(random_state=42))])
```

```

pipeline_dt=Pipeline([
    ↪('dt_classifier',DecisionTreeClassifier(random_state=42))]

pipeline_rf=Pipeline([('rf_classifier',RandomForestClassifier())])

pipeline_svc=Pipeline([('sv_classifier',SVC())])

pipeline_knn=Pipeline([('KNN_Classifier',KNeighborsClassifier())])

# List of all the pipelines

pipelines = [pipeline_lr, pipeline_dt, pipeline_rf, pipeline_svc, pipeline_knn]

# Dictionary of pipelines and classifier types for ease of reference

pipe_dict = {0: 'Logistic Regression', 1: 'Decision Tree', 2: 'RandomForest', 3:
    ↪ "SVC", 4:'Knn Neighbours'}

# Fit the pipelines

# for pipe in pipelines:
#     pipe.fit(X_train, Y_train)

```

```
[ ]: df['GENDER'].unique()
```

```
[ ]: array([0, 1], dtype=int64)
```

```
[ ]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 309 entries, 0 to 308
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   GENDER                309 non-null    int64
1   AGE                   309 non-null    int64
2   SMOKING                309 non-null    int64
3   YELLOW_FINGERS        309 non-null    int64
4   ANXIETY                309 non-null    int64
5   PEER_PRESSURE         309 non-null    int64
6   CHRONIC_DISEASE       309 non-null    int64
7   FATIGUE                309 non-null    int64
8   ALLERGY                309 non-null    int64
9   WHEEZING              309 non-null    int64
10  ALCOHOL_CONSUMING     309 non-null    int64
11  COUGHING               309 non-null    int64

```



```

12  SHORTNESS OF BREATH      309 non-null    int64
13  SWALLOWING DIFFICULTY   309 non-null    int64
14  CHEST PAIN               309 non-null    int64
15  LUNG_CANCER              309 non-null    int64
dtypes: int64(16)
memory usage: 38.8 KB

```

```

[ ]: # Fit the pipelines

for pipe in pipelines:
    pipe.fit(X_train, Y_train)

```

```

[ ]: #cross validation on accuracy

cv_results_accuracy = []
for i, model in enumerate(pipelines):
    cv_score = cross_val_score(model, X_train,Y_train, cv=10 )
    cv_results_accuracy.append(cv_score)
    print("%s: %f " % (pipe_dict[i], cv_score.mean()))

```

```

Logistic Regression: 0.914833
Decision Tree: 0.858500
RandomForest: 0.906833
SVC: 0.874500
Knn Neighbours: 0.874500

```

```

[ ]: # LogisticRegression gave more accuracy score so we choose LR

lr = LogisticRegression()
lr.fit(X_train, Y_train)
y_pred = lr.predict(X_test)

lr_train_acc = accuracy_score(Y_train, lr.predict(X_train))
lr_test_acc = accuracy_score(Y_test, y_pred)
cm = confusion_matrix(Y_test, y_pred)

print(f"Training Accuracy of Logistic Regression Model is {lr_train_acc}")
print(f"Test Accuracy of Logistic Regression Model is {lr_test_acc}")

print(f"confusion_matrix is\n {cm}")

```

```

Training Accuracy of Logistic Regression Model is 0.9392712550607287
Test Accuracy of Logistic Regression Model is 0.9193548387096774
confusion_matrix is
[[53  1]
 [ 4  4]]

```

```

[ ]: print(classification_report(Y_test, y_pred))

```

	precision	recall	f1-score	support
0	0.93	0.98	0.95	54
1	0.80	0.50	0.62	8
accuracy			0.92	62
macro avg	0.86	0.74	0.79	62
weighted avg	0.91	0.92	0.91	62